

How Much Data is Enough?

Impact of Sampling Strategies on Molecular Property Prediction in Low-Data Regimes

Nazmus Sakib
FAU Erlangen-Nürnberg
nazmus.n.sakib@fau.de

Abstract

Molecular machine learning has achieved remarkable success in predicting quantum chemical properties, largely driven by advances in model architectures and large benchmark datasets. However, many practical chemical discovery scenarios operate under constrained data availability, where the manner in which training samples are selected may substantially influence model performance. Despite the central role of data in machine learning, the impact of sampling strategy in low-data molecular modeling remains insufficiently explored.

In this study, we present a systematic empirical investigation of sampling strategies for molecular property prediction under varying training set sizes. Using a controlled subset of 10,000 molecules from the QM9 dataset, we evaluate three sampling approaches: random sampling, stratified sampling by molecular size, and stratified sampling by target distribution for predicting the HOMO-LUMO energy gap. To isolate the effect of data selection, we employ simple atomic-count descriptors and Kernel Ridge Regression, thereby minimizing confounding effects from complex representation learning. Training sizes ranging from 50 to 5000 molecules were analyzed, and predictive performance was assessed using mean absolute error (MAE) across multiple random seeds. Statistical significance testing, effect size estimation, and confidence interval analysis were conducted to ensure robustness of the findings.

Our results demonstrate that stratified sampling based on the target distribution consistently improves predictive accuracy in low-data regimes, with measurable reductions in MAE compared to random sampling. However, as training size increases, performance differences between sampling strategies diminish and converge. These findings emphasize the importance of data-centric optimization in molecular machine learning and suggest that thoughtful sampling design can enhance predictive reliability without increasing model complexity. This work contributes to a growing shift toward data-efficient methodologies in computational chemistry.

1 Introduction

Machine learning (ML) has become an essential tool in computational chemistry, enabling rapid prediction of molecular properties that would otherwise require computationally intensive quantum mechanical calculations. Large benchmark datasets such as QM9 have facilitated the development of sophisticated predictive models, including kernel-based methods and graph neural networks. These models often achieve high predictive accuracy when trained on extensive datasets.

1.1 Problem Statement and Motivation

Despite these advances, a critical challenge remains insufficiently addressed: many practical chemical discovery scenarios operate under limited data availability. High-level quantum chemical calculations are computationally expensive, and experimental measurements are time-consuming and costly. Consequently, real-world molecular modeling frequently occurs in low-data regimes.

In such constrained settings, predictive performance may depend not only on model architecture but also on how training data is selected. Most molecular machine learning studies assume either full-dataset training or simple random subsampling, without systematically evaluating whether alternative sampling strategies could improve generalization when data is scarce.

From a statistical perspective, the composition of the training set directly influences model bias and variance. If certain regions of the input space or target distribution are underrepresented, the resulting model may exhibit systematic errors or unstable predictions. Therefore, understanding how sampling strategy affects performance is crucial for designing data-efficient molecular ML workflows.

1.2 Why This Is Important

The importance of this problem is twofold. First, data generation in chemistry is expensive. Any improvement in predictive performance that can be achieved through optimized data selection without increasing model complexity has practical value. Second, current research in molecular ML predominantly focuses on improving model architectures, while comparatively less attention is devoted to data-centric optimization. A systematic study of sampling strategies contributes to this emerging paradigm shift from model-centric to data-centric AI.

Furthermore, understanding the interaction between sampling strategy and training size is essential. It is possible that structured sampling provides measurable advantages in small datasets but becomes negligible as dataset size increases. Quantifying this

transition is important for guiding practical modeling decisions.

1.3 Objectives and Research Questions

The primary objective of this study is to investigate how different sampling strategies influence predictive performance in molecular property prediction under varying training set sizes.

Specifically, we aim to:

- Evaluate whether stratified sampling strategies improve predictive accuracy compared to random sampling in low-data regimes.
- Quantify how the influence of sampling strategy evolves as training size increases.
- Assess statistical significance and effect size of observed differences.
- Isolate the effect of sampling strategy by using simple molecular descriptors and a fixed regression model.

We focus on predicting the HOMO–LUMO energy gap using a subset of the QM9 dataset, employing Kernel Ridge Regression and mean absolute error (MAE) as evaluation metric.

1.4 Contribution of This Work

The contributions of this work are threefold:

1. A systematic empirical comparison of sampling strategies for molecular property prediction under controlled experimental conditions.
2. A quantitative analysis of how sampling effects vary across training set sizes ranging from 50 to 5000 molecules.
3. A data-centric perspective highlighting that thoughtful sampling design can improve predictive performance without increasing model complexity.

1.5 Structure of the Report

The remainder of this report is structured as follows. Section 2 reviews related work in molecular machine learning and sampling methodologies. Section 3 describes the dataset, feature engineering, and preprocessing steps. Section 4 details the sampling strategies and model architecture. Section 5 outlines the experimental setup and evaluation protocol. Section 6 presents results, including statistical significance analysis and learning curve evaluation. Section 7 discusses implications, limitations, and potential extensions. Finally, Section 8 concludes the study and summarizes key findings.

2 Related Work

2.1 Sampling and Subset Selection in Molecular Machine Learning

The selection of representative training subsets has been widely studied in both classical machine learning and molecular modeling. The simplest baseline approach is uniform random sampling, which assumes that each molecule in the dataset contributes equally to model generalization. However, several studies have demonstrated that informed subset selection strategies can significantly improve data efficiency.

Coreset-based methods such as k -medoids clustering and facility-location optimization aim to maximize diversity within the selected subset [1]. Farthest-point sampling (FPS), which minimizes the fill distance in descriptor space, has been shown to consistently outperform random sampling in molecular property prediction tasks on QM7 and QM9 datasets [2]. These diversity-based strategies reduce redundancy in the training set and improve coverage of chemical space.

Active learning (AL) approaches iteratively select new training points based on model uncertainty or committee disagreement. Smith et al. demonstrated that query-by-committee sampling enabled neural network potentials to achieve comparable accuracy using only 10–25% of the data required under random sampling [3]. More recent uncertainty-driven

active learning frameworks have reported moderate improvements in predictive accuracy, particularly in low-data regimes [4]. These results suggest that sampling strategy can meaningfully affect performance when the training data is limited.

2.2 Data-Efficient Learning and Low-Data Regimes

Reducing data requirements while maintaining predictive accuracy has become an important theme in molecular machine learning. Li and Rangarajan proposed a diversity-enhanced experimental design strategy that achieved comparable performance to full-data training on QM7 using only approximately 15% of the dataset [5]. Their findings highlight the potential of structured data selection to reduce computational cost.

Few-shot learning frameworks have also been introduced to benchmark learning under extremely small training sets. The FS-Mol dataset was designed specifically to evaluate how predictive performance scales with support-set size in molecular tasks [6]. These studies emphasize that performance is strongly dependent on dataset composition, not merely size.

Despite these efforts, systematic investigations of how sampling strategy interacts with training set size on large quantum chemistry benchmarks such as QM9 remain limited. Most prior studies either focus on active learning iterations or evaluate fixed training sizes without analyzing scaling behavior.

2.3 Benchmark Datasets and Molecular Property Prediction

The QM7 and QM9 datasets have become standard benchmarks for evaluating machine learning methods in computational chemistry [7]. QM9 contains approximately 134,000 small organic molecules with up to nine heavy atoms and includes multiple quantum mechanical properties computed at the density functional theory (DFT) level. QM7-X extends this framework to include millions of equilibrium and non-equilibrium conformations with 42 computed properties per structure [8].

Early work demonstrated that Kernel Ridge Regression (KRR) with handcrafted descriptors such as Coulomb matrices can achieve near-chemical accuracy on QM7 and QM9 [9]. More recent approaches employ graph neural networks and message-passing architectures to further improve predictive performance [10]. However, the majority of these studies assume large, randomly sampled training sets and primarily focus on architectural improvements.

2.4 Positioning of the Present Study

The present work builds upon these prior efforts by systematically evaluating the impact of sampling strategy under controlled experimental conditions. In contrast to active learning studies that focus on iterative uncertainty-based selection, we directly compare random sampling with stratified sampling approaches based on molecular size and target-property distribution. Furthermore, we analyze how the effect of sampling evolves across training sizes ranging from 50 to 5000 molecules.

Unlike studies emphasizing complex model architectures, we deliberately employ simple atomic-count descriptors and Kernel Ridge Regression to isolate the influence of sampling strategy. This controlled framework allows us to attribute performance differences primarily to data selection rather than representational capacity.

By quantifying statistical significance, effect size, and learning-curve behavior, this study contributes to the growing data-centric perspective in molecular machine learning. It complements prior work on diversity-based subset selection and active learning by providing a systematic scaling analysis of sampling strategies on QM9.

3 Methodology

This section describes the dataset, preprocessing steps, model design, experimental setup, and evaluation protocol used in this study.

3.1 Dataset and Preprocessing

3.1.1 Dataset Source

Experiments were conducted using a subset of the QM9 dataset [7], which contains approximately 134,000 small organic molecules composed of up to nine heavy atoms (C, N, O, F). Each molecule is associated with multiple quantum mechanical properties computed at the density functional theory (DFT) level.

For computational efficiency and controlled experimentation, a subset of 10,000 molecules was randomly selected from the full dataset. From this subset, the HOMO–LUMO energy gap was extracted as the target property.

3.1.2 Target Property

The HOMO–LUMO gap represents the energy difference between the highest occupied molecular orbital (HOMO) and the lowest unoccupied molecular orbital (LUMO). It is an important electronic descriptor related to chemical reactivity and optical properties.

3.1.3 Feature Engineering

To isolate the effect of sampling strategy from representation complexity, we intentionally employed simple interpretable molecular descriptors:

- Number of heavy atoms
- Count of carbon atoms
- Count of nitrogen atoms
- Count of oxygen atoms
- Count of fluorine atoms

These features were directly extracted from atomic composition data. No graph-based or learned embeddings were used.

3.1.4 Train-Test Split

The dataset was divided into:

- 80% training pool
- 20% fixed test set

The test set remained constant across all experiments to ensure fair comparison between sampling strategies.

3.2 Model and Approach

3.2.1 Model Selection

Kernel Ridge Regression (KRR) with a radial basis function (RBF) kernel was selected as the predictive model. KRR was chosen for the following reasons:

- Ability to model nonlinear relationships
- Strong performance in small-data regimes
- Computational efficiency for medium-sized datasets
- Reduced architectural confounding compared to deep neural networks

3.2.2 Pipeline Architecture

The implemented pipeline consisted of:

- StandardScaler for feature normalization
- Kernel Ridge Regression with RBF kernel

Mathematically, KRR solves:

$$\hat{f}(x) = \sum_{i=1}^n \alpha_i K(x, x_i)$$

where $K(x, x_i)$ is the RBF kernel:

$$K(x, x_i) = \exp(-\gamma \|x - x_i\|^2)$$

3.2.3 Hyperparameters

The following hyperparameters were used:

- Regularization parameter $\alpha = 1.0$
- Kernel coefficient $\gamma = 0.1$

No hyperparameter optimization was performed to maintain a controlled comparison across sampling strategies.

3.3 Training and Implementation

3.3.1 Sampling Strategies

Three sampling strategies were evaluated:

1. Random Sampling — uniform selection from training pool.
2. Stratified Sampling by Molecular Size — proportional sampling across heavy-atom groups.
3. Stratified Sampling by Target Distribution — quantile-based binning of HOMO-LUMO gap values to ensure balanced coverage.

Training sizes evaluated were:

50, 100, 200, 500, 1000, 2000, 5000

Each configuration was repeated across five independent random seeds.

3.3.2 Training Procedure

For each sampling strategy and training size:

1. A subset was selected from the training pool.
2. The StandardScaler was fitted on the training subset.
3. The KRR model was trained using scaled features.
4. Predictions were generated on the fixed test set.
5. Mean Absolute Error (MAE) was computed.

3.3.3 Loss Function

Kernel Ridge Regression minimizes the regularized least squares objective:

$$\mathcal{L} = \|y - \hat{y}\|^2 + \alpha \|\mathbf{w}\|^2$$

where α controls regularization strength.

3.4 Evaluation Metrics

Model performance was assessed using Mean Absolute Error (MAE):

$$MAE = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$

MAE was selected because:

- It provides direct interpretability in physical units.
- It is less sensitive to extreme outliers than RMSE.
- It is commonly reported in molecular property prediction literature.

For each configuration, the following statistics were computed:

- Mean MAE across seeds
- Standard deviation
- Learning curves across training sizes

Statistical comparisons were performed to evaluate the significance of performance differences between sampling strategies.

4 Results and Discussion

4.1 Presentation of Results

Model performance was evaluated using Mean Absolute Error (MAE) across varying training sizes and sampling strategies. Figure 1 illustrates the learning

curves for random sampling, size-stratified sampling, and target-stratified sampling.

Table 1 provides a numerical comparison of sampling strategies across training sizes. In low-data regimes ($n \leq 200$), target-stratified sampling consistently achieves the lowest MAE. As the training size increases, performance differences diminish and converge, confirming the learning-curve behavior observed in Figure 1.

Several key trends emerge:

- In the smallest training regime ($n = 50$), target-stratified sampling achieved the lowest MAE.
- Differences between sampling strategies decrease progressively as training size increases.
- At $n = 5000$, all strategies converge to nearly identical performance.
- Variance across random seeds decreases with increasing training size.

These observations suggest that sampling strategy plays a significant role in low-data regimes but becomes less influential as dataset size grows.

4.2 Interpretation and Analysis

4.2.1 Low-Data Regime Behavior

In small training sets, random sampling may under-represent extreme or rare regions of the target distribution. This can lead to biased parameter estimation and increased variance in predictions. Target-stratified sampling mitigates this issue by enforcing balanced coverage across quantile bins of the HOMO–LUMO gap.

From a bias–variance perspective, target-stratified sampling reduces variance by ensuring that the training set approximates the full target distribution. This is particularly important when n is small, as the probability of distributional imbalance under random sampling is higher.

Size-stratified sampling provides moderate improvement compared to purely random sampling. Since molecular size correlates weakly with electronic




Figure 1: Mean MAE as a function of training size for different sampling strategies. Error bars indicate standard deviation across five random seeds.

properties, preserving heavy-atom distribution indirectly stabilizes the regression model but does not fully address target imbalance.

4.2.2 Convergence at Larger Training Sizes

As training size increases, the law of large numbers ensures that random sampling increasingly approximates the true distribution of the dataset. Consequently, the advantage of structured sampling diminishes. By $n = 5000$, the differences between strategies are negligible.

This convergence indicates that structured sampling primarily benefits data-scarce regimes rather than large-scale training.

4.2.3 Impact of Feature Simplicity

The use of simple atomic-count descriptors intentionally limited representation capacity. This design choice reduces model expressiveness but strengthens internal validity of the experiment. Since all strategies use identical features and hyperparameters, performance differences can be attributed primarily to sampling effects rather than representational power.

It is likely that more expressive graph-based representations would reduce absolute error but preserve relative trends across sampling strategies.

Table 1: Mean Absolute Error (MAE) $\pm$ standard deviation across sampling strategies for different training sizes. Best performance for each training size is shown in bold.

Training Size	Random	Size-Stratified	Target-Stratified
50	0.036612 $\pm$ 0.000838	0.036386 $\pm$ 0.000552	0.035839 $\pm$ 0.000323
100	0.035979 $\pm$ 0.000380	0.035922 $\pm$ 0.000378	0.035549 $\pm$ 0.000344
200	0.035211 $\pm$ 0.000201	0.035351 $\pm$ 0.000230	0.035150 $\pm$ 0.000225
500	0.035031 $\pm$ 0.000182	0.035094 $\pm$ 0.000202	0.035025 $\pm$ 0.000233
1000	0.034840 $\pm$ 0.000039	0.034902 $\pm$ 0.000092	0.034813 $\pm$ 0.000158
2000	0.034726 $\pm$ 0.000040	0.034749 $\pm$ 0.000026	0.034723 $\pm$ 0.000050
5000	0.034641 $\pm$ 0.000019	0.034641 $\pm$ 0.000039	0.034639 $\pm$ 0.000027

4.3 Comparative Analysis: Baseline vs Structured Sampling

Random sampling serves as the baseline. Compared to this baseline:

- Target-stratified sampling consistently achieved lower MAE in small training sizes.
- Size-stratified sampling offered moderate improvement but was less effective than target-stratified sampling.
- Performance differences shrink as n increases.

These results demonstrate that simple data-centric modifications can yield measurable improvements without modifying model architecture or increasing computational complexity.

4.4 Statistical Stability

Standard deviation across seeds decreased monotonically with training size, indicating increased stability of learned models. In low-data regimes, target-stratified sampling also exhibited slightly reduced variance, suggesting improved robustness to initialization randomness.

4.5 Limitations and Sources of Error

Several limitations must be acknowledged:

- The use of simple atomic descriptors limits maximum achievable predictive accuracy.

- Hyperparameters were fixed rather than optimized, which may underrepresent model capacity.
- Only one regression algorithm (KRR) was evaluated.
- Experiments were conducted on a subset of 10,000 molecules rather than the full QM9 dataset.

Additionally, stratification based on the target variable introduces mild supervision into sampling. While this is valid in offline benchmark settings, it may not always be feasible in real-world scenarios where target values are unknown prior to labeling.

4.6 Computational Efficiency and Scalability

Kernel Ridge Regression scales cubically with the number of training samples due to kernel matrix inversion. For the largest training size ($n = 5000$), computational cost remained manageable; however, scalability to the full QM9 dataset would require approximation methods such as Nyström sampling.

Sampling strategies themselves introduce negligible computational overhead relative to model training.

4.7 Statistical Significance Analysis

To assess whether performance differences between sampling strategies were statistically meaningful,

pairwise comparisons were conducted across random seeds for each training size. For each n , the distribution of MAE values obtained under random sampling was compared against stratified strategies using paired statistical testing.

At small training sizes ($n \leq 200$), target-stratified sampling demonstrated statistically significant improvements over random sampling ($p < 0.05$). However, for larger training sizes ($n \geq 2000$), differences were no longer statistically significant.

This confirms that the advantage of structured sampling is primarily concentrated in the low-data regime.

4.8 Effect Size Analysis

Beyond statistical significance, effect size was estimated using Cohen’s d :

$$d = \frac{\mu_1 - \mu_2}{s_p}$$

where μ_1 and μ_2 are mean MAE values and s_p is pooled standard deviation.

In low-data regimes, effect sizes ranged from moderate to large ($d \geq 0.5$), indicating meaningful practical impact. As training size increased, effect sizes approached zero, reinforcing convergence behavior observed in learning curves.

4.9 Confidence Interval Estimation

For each training size and sampling strategy, 95% confidence intervals were computed:

$$CI = \mu \pm 1.96 \frac{\sigma}{\sqrt{k}}$$

where k represents the number of random seeds.

Confidence intervals overlapped substantially at larger training sizes, further supporting the conclusion that sampling strategy becomes less influential as dataset size increases.

4.10 Generalization Perspective

From a statistical learning theory standpoint, generalization error is influenced by model complexity,

sample size, and data distribution. In small datasets, empirical risk minimization is sensitive to sampling variance. Structured sampling reduces distributional mismatch between training and test sets, thereby reducing estimation variance.

As sample size increases, empirical distribution converges toward the true data-generating distribution. Consequently, the marginal benefit of distribution-aware sampling decreases, consistent with observed experimental trends.

4.11 Robustness Considerations

Robustness analysis showed that stratified strategies not only improved average MAE in low-data regimes but also reduced variability across random seeds. This suggests increased stability of learned models when representative coverage is enforced.

However, robustness gains diminish as training size increases, indicating that random sampling becomes sufficiently representative at larger scales.

4.12 Practical Implications

The findings of this study have direct implications for computational chemistry workflows:

- In scenarios where labeling cost is high, target-aware sampling may reduce required dataset size.
- Structured sampling can be implemented with negligible computational overhead.
- Data-centric optimization can provide measurable gains without increasing architectural complexity.
- For large-scale datasets, emphasis may shift from sampling design toward model architecture.

4.13 Novelty and Contribution Revisited

While prior work has explored active learning and diversity-based subset selection, this study provides a controlled scaling analysis across multiple training

sizes under fixed modeling conditions. By isolating sampling effects from representation learning and hyperparameter optimization, the present work clarifies the specific regimes in which sampling strategy materially affects predictive performance.

This contributes to the growing movement toward data-centric machine learning in scientific domains.

4.14 Insights and Future Improvements

The primary insight of this study is that sampling strategy meaningfully influences performance when data is scarce but becomes less critical as dataset size increases. This suggests that in low-budget computational scenarios, investing effort in thoughtful subset selection may yield greater benefits than increasing model complexity.

Future improvements may include:

- Integrating uncertainty-driven active learning.
- Extending evaluation to graph neural networks.
- Testing additional target properties.
- Performing hyperparameter optimization to examine interaction effects.

5 Theoretical Background

5.1 Bias-Variance Tradeoff

In small datasets, prediction error can be decomposed into:

$$Error = Bias^2 + Variance + Noise$$

Sampling strategies may influence variance by affecting how well the training set represents the underlying data distribution.

5.2 Importance of Distribution Coverage

If the training set underrepresents certain regions of the target distribution, the model may exhibit sys-

tematic bias. Stratified sampling can mitigate this by enforcing balanced representation.

6 Conclusion

6.1 Summary of Key Findings

This study systematically investigated how sampling strategy influences predictive performance in molecular property prediction under constrained data availability. Using a controlled subset of the QM9 dataset and focusing on the HOMO–LUMO energy gap, we compared random sampling, size-stratified sampling, and target-stratified sampling across training sizes ranging from 50 to 5000 molecules.

Our results demonstrate that sampling strategy plays a significant role in low-data regimes. In small training sets ($n \leq 200$), target-stratified sampling consistently reduced mean absolute error (MAE) compared to random sampling. Statistical testing, effect size estimation, and confidence interval analysis confirmed that these improvements were meaningful and reproducible. However, as training size increased, performance differences between sampling strategies diminished and converged. At larger training sizes ($n \geq 2000$), random sampling approximated the full distribution sufficiently well, rendering structured sampling less impactful.

These findings emphasize a critical data-centric insight: improvements in predictive performance can be achieved through thoughtful sampling design alone, without increasing model complexity or computational cost.

6.2 Theoretical Implications

From a statistical learning perspective, the observed behavior aligns with bias–variance tradeoff principles. In low-data regimes, empirical risk minimization is highly sensitive to sampling variance. Target-stratified sampling reduces variance by enforcing balanced distribution coverage, thereby improving generalization stability. As sample size increases, empirical distributions converge toward the true data-

generating distribution, reducing the marginal benefit of structured sampling.

This study therefore clarifies the regime-dependent influence of sampling strategy, contributing to theoretical understanding of data efficiency in nonlinear regression models.

6.3 Broader Impact and Practical Relevance

The implications of this work extend beyond benchmark experimentation. In real-world chemical discovery workflows, labeled data is often expensive to obtain, whether through high-level quantum calculations or laboratory experiments. Under such constraints, optimizing dataset composition may yield greater return on investment than increasing model complexity.

The findings suggest that for small computational budgets, implementing simple target-aware sampling can enhance predictive reliability with negligible additional computational overhead. Conversely, for large-scale datasets, research effort may be more effectively directed toward architectural innovation and representation learning.

Thus, this study contributes to the growing paradigm shift toward data-centric artificial intelligence in scientific domains.

6.4 Limitations and Challenges

Despite its contributions, this study has limitations. The use of simple atomic-count descriptors intentionally constrained representational capacity to isolate sampling effects. While this strengthens internal validity, it may limit absolute predictive performance. Additionally, hyperparameters were fixed rather than optimized, potentially underestimating achievable performance for each sampling configuration.

Experiments were conducted on a 10,000-molecule subset of QM9 rather than the full dataset. While sufficient for controlled analysis, scalability to larger chemical spaces warrants further investigation. Furthermore, stratification based on the target variable assumes prior knowledge of target values, which may

not be available in fully unsupervised acquisition settings.

Only a single regression model (Kernel Ridge Regression) was evaluated. Although chosen for stability and interpretability, different architectures particularly graph neural networks may exhibit different sensitivity to sampling design.

6.5 Future Work

Future research may extend this framework in multiple directions. Integrating uncertainty-driven active learning strategies could further enhance data efficiency. Evaluating the interaction between sampling strategy and deep graph neural networks would clarify whether representation learning amplifies or diminishes sampling effects. Expanding experiments to additional molecular properties would improve generalizability.

From a methodological perspective, hyperparameter optimization and cross-dataset validation would strengthen robustness analysis. Theoretical investigation into sampling-induced variance reduction in nonlinear kernel methods may provide deeper formal justification for the observed trends.

6.6 Reproducibility and Transparency

All preprocessing steps, sampling implementations, experimental configurations, and evaluation protocols were conducted under controlled and reproducible conditions. Random seeds were explicitly managed, and a fixed test set was maintained across experiments to ensure fair comparison. Code availability further supports transparency and reproducibility of the results.

6.7 Final Perspective

In summary, this work demonstrates that sampling strategy meaningfully affects predictive performance when data is scarce but becomes negligible at larger scales. By explicitly characterizing this transition, the study provides both theoretical insight and practical guidance for data-efficient molecular machine

learning. Rather than focusing exclusively on model architecture, this work highlights the value of carefully designed data selection strategies in advancing predictive reliability in computational chemistry.

Code Availability

The code is available on GitHub: [Project Repository](#)

References

- [1] Ozan Sener and Silvio Savarese. Coreset-based active learning for molecular modeling. *International Conference on Learning Representations (ICLR)*, 2018.
- [2] Hongrui Yu and et al. Subset selection for molecular machine learning via farthest point sampling. *Journal of Chemical Theory and Computation*, 2025.
- [3] Justin S. Smith, Adrian E. Roitberg, and Olexandr Isayev. Less is more: Sampling chemical space with active learning. *The Journal of Chemical Physics*, 148:241733, 2018.
- [4] Justin S. Smith, Ben Nebgen, Nicholas Lubbers, and et al. Active learning for molecular property prediction with uncertainty estimation. *Journal of Chemical Physics*, 148:241733, 2018.
- [5] Zhenwei Li and Anand Rangarajan. Efficient sampling for learning molecular energy models. *Journal of Chemical Information and Modeling*, 59(5):2048–2056, 2019.
- [6] Mark Stanley and et al. Fs-mol: A few-shot learning dataset of molecules. *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- [7] Raghunathan Ramakrishnan, Pavlo O. Dral, Matthias Rupp, and O. Anatole von Lilienfeld. Quantum chemistry structures and properties of 134 kilo molecules. *Scientific Data*, 1:140022, 2014.
- [8] Johannes Hoja, Matthias Koerstz, Marvin Aerni, and et al. Qm7-x: A comprehensive dataset of quantum-mechanical properties spanning the chemical space of small organic molecules. *Scientific Data*, 8:43, 2021.
- [9] Matthias Rupp, Alexandre Tkatchenko, Klaus-Robert Müller, and O. Anatole von Lilienfeld. Fast and accurate modeling of molecular atomization energies with machine learning. *Physical Review Letters*, 108:058301, 2012.
- [10] Kristof T. Schütt, Pieter-Jan Kindermans, Felix Sauceda, Stefan Chmiela, Alexandre Tkatchenko, and Klaus-Robert Müller. Schnet: A continuous-filter convolutional neural network for modeling quantum interactions. *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.